

# CSP554—Big Data Technologies

## Assignment #9

**Worth: 15 points (1 point for each problem)**

## HBase

Start up a Hadoop cluster as previously, but instead of choosing the “Core Hadoop” configuration chose the “HBase” configuration (see below), otherwise proceed as before:

[illegible]

## Exercises

Exercise 1) (1 point)

## Create an HBase table with the following characteristics

Table Name: csp554Tbl

First column family: cf1

Second column family: cf2

Then execute the DESCRIBE command on the table and return command you wrote and the output as the results of this exercise.

**Command:**

```
Hbase shell
```

```
create 'csp554Tbl', {NAME=>'cf1'}, {NAME=>'cf2'}
```

```
describe 'csp554Tb1'
```

```
hbase(main):001:0> create 'csp554tb1', {NAME=>'cf1'}, {NAME=>'cf2'}
0 row(s) in 1.8150 seconds

=> Hbase::Table - csp554tb1
hbase(main):002:0> describe 'csp554tb1'
table csp554tb1 is ENABLED
csp554tb1
COLUMN FAMILIES DESCRIPTION
{NAME => 'cf1', BLOOMFILTER => 'Row', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
{NAME => 'cf2', BLOOMFILTER => 'Row', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
2 row(s) in 0.0470 seconds

hbase(main):003:0>
```

### Exercise 2) (1 point)

Put the following data into the table created in exercise 1:

Row Key	Column Family	Column (Qualifier)	Value
Row1	cf1	name	Sam
Row2	cf1	name	Ahmed
Row1	cf2	job	Pilot
Row2	cf2	job	Doctor
Row1	cf2	level	LZ3
Row2	cf2	level	AR7

Execute the SCAN command on this table returning all rows, column families and columns. Provide the command and its result as the output of this exercise.

#### Command:

```
put 'csp554Tbl', 'Row1', 'cf1:name', 'Sam'
put 'csp554Tbl', 'Row2', 'cf1:name', 'Ahmed'
put 'csp554Tbl', 'Row1', 'cf2:job', 'Pilot'
put 'csp554Tbl', 'Row2', 'cf2:job', 'Doctor'
put 'csp554Tbl', 'Row1', 'cf2:level', 'LZ3'
put 'csp554Tbl', 'Row2', 'cf2:level', 'AR7'
scan 'csp554Tbl'
```

```
hbase(main):009:0> scan 'csp554Tbl'
ROW                                COLUMN+CELL
Row1                               column=cf1:name, timestamp=1669830513131, value=Sam
Row1                               column=cf2:job, timestamp=1669830538538, value=Pilot
Row1                               column=cf2:level, timestamp=1669830555047, value=LZ3
Row2                               column=cf1:name, timestamp=1669830530721, value=Ahmed
Row2                               column=cf2:job, timestamp=1669830546275, value=Doctor
Row2                               column=cf2:level, timestamp=1669830563575, value=AR7
2 row(s) in 0.0360 seconds
hbase(main):010:0> |
```

### Exercise 3) (1 point)

Using the above table write a command that will get the value associated with row (Row1), column family (cf2) and column/qualifier (level). Provide the command and its result as the output of this exercise.

#### Command:

```
get 'csp554Tbl', 'Row1', 'cf2:level'
```

```
hbase(main):010:0> get 'csp554Tbl','Row1','cf2:level'
COLUMN                                CELL
cf2:level                             timestamp=1669830555047, value=LZ3
1 row(s) in 0.0410 seconds
hbase(main):011:0> |
```

#### Exercise 4) (1 point)

Using the above table write command that will get the value associated with row (Row2), column family (cf1) and column/qualifier (name). Provide the command and its result as the output of this exercise.

##### **Command:**

```
get 'csp554Tbl','Row2','cf1:name'
```

```
hbase(main):011:0> get 'csp554Tbl','Row2','cf1:name'
COLUMN                                CELL
cf1:name                             timestamp=1669830530721, value=Ahmed
1 row(s) in 0.0260 seconds
hbase(main):012:0> |
```

#### Exercise 5) (1 point)

Using the above table write a SCAN command that will return information about only two rows using the LIMIT modifier. Provide the command and its result as the output of this exercise.

##### **Command:**

```
scan 'csp554Tbl',{LIMIT=>1}
```

```
hbase(main):012:0> scan 'csp554Tbl',{LIMIT=>1}
ROW                                  COLUMN+CELL
Row1                                 column=cf1:name, timestamp=1669830513131, value=Sam
Row1                                 column=cf2:job, timestamp=1669830538538, value=Pilot
Row1                                 column=cf2:level, timestamp=1669830555047, value=LZ3
1 row(s) in 0.0380 seconds
hbase(main):013:0> |
```

## Cassandra

### Exercises:

#### Exercise 1) (1 point)

Read the article “A Big Data Modeling Methodology for Apache Cassandra” and provide a ½ page summary including your comments and impressions.

Leading distributed database that is highly available, scalable, and transactional is Apache Cassandra. On clusters with thousands of nodes spread across numerous data centers, it is recognized to manage some of the largest databases in the world. Product catalogs and playlists, sensor data and the Internet of Things, messaging and social networking, recommendation, personalization, fraud detection, and several more time series data-related applications are examples of Cassandra data management use cases. The widespread use of Cassandra in big data applications is due, in part, to its scalable and fault-tolerant peer-to-peer architecture, adaptable and flexible data model that developed from the BigTable data model, declarative and user-friendly Cassandra Query Language (CQL), and extremely efficient write and read access paths that enable crucial big data applications to remain always available.

In this paper, a rigorous query-driven data modeling approach for Apache Cassandra is presented. In a variety of areas, including query-driven schema design, data nesting, and data duplication, the methodology was demonstrated to be significantly different from the conventional relational data modelling approach. In order to move from conceptual data models that are independent of technology to logical data models that are specific to Cassandra, this article describes mapping rules and mapping patterns and elaborates on the core data modeling concepts for Cassandra. Additionally, it clarifies the function of physical data modeling and suggests Chebotko Diagrams, a cutting-edge visualization method that may be utilized to record intricate logical and physical data models.

Finally, it introduces KDM, a potent data modeling tool that automates some of the most difficult, error-prone, and time-consuming data modelling processes, such as conceptual-to-logical mapping, logical-to-physical mapping, and CQL creation.

#### Exercise 2) (1 point)

##### Step A – Start an EMR cluster

Start up an EMR/Hadoop cluster as previously, but instead of choosing the “Core Hadoop” configuration chose the “Spark” configuration (see below), otherwise proceed as before.

##### Step B – Install the Cassandra database software and start it

Open up a terminal connection to your EMR master node. Over the course of this exercise, you will need to open up three separate terminal connections to your EMR master node. This is the first, which we will call Cass-Term.







c) To check if your script file has created a keyspace execute the following in the CQL shell:

describe keyspaces;

```
cqlsh> describe keyspaces;  
system_schema system_auth a20506653 system system_distributed system_traces  
cqlsh> |
```

d) At this point you have created a keyspace unique to you. So, make that keyspace the default by entering the following into the CQL shell:

USE A20506653;

```
cqlsh> USE A20506653;  
cqlsh:a20506653> |
```

Now create a file in your working directory called ex2.cql using the Edit-Term (or PC/MAC and scp). In this file write the command to create a table named 'Music' with the following characteristics:

Attribute Name	Attribute Type	Primary Key / Cluster Key
<b>artistName</b>	text	Primary Key
<b>albumName</b>	text	Cluster Key
<b>numberSold</b>	int	Non Key Column
<b>Cost</b>	int	Non Key Column

Execute ex2.cql in the CQL shell. Then execute the shell command 'DESCRIBE TABLE Music;' and include the output as the result of this exercise.

```
[hadoop@ip-172-31-44-252 ~]$ cat ex2.cql  
CREATE table a20506653.music(  
  artistname text,  
  albumname text,  
  numbersold int,  
  cost int,  
  PRIMARY KEY(artistname, albumname)  
) WITH CLUSTERING ORDER BY (albumname DESC);  
[hadoop@ip-172-31-44-252 ~]$ |
```

```

cqlsh:a20506653> describe table music;

CREATE TABLE a20506653.music (
  artistname text,
  albumname text,
  cost int,
  numbersold int,
  PRIMARY KEY (artistname, albumname)
) WITH CLUSTERING ORDER BY (albumname DESC)
AND bloom_filter_fp_chance = 0.01
AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
AND comment = ''
AND compaction = {'class': 'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy', 'max_threshold': '32', 'min_threshold': '4'}
AND compression = {'chunk_length_in_kb': '64', 'class': 'org.apache.cassandra.io.compress.LZ4Compressor'}
AND crc_check_chance = 1.0
AND dclocal_read_repair_chance = 0.1
AND default_time_to_live = 0
AND gc_grace_seconds = 864000
AND max_index_interval = 2048
AND memtable_flush_period_in_ms = 0
AND min_index_interval = 128
AND read_repair_chance = 0.0
AND speculative_retry = '99PERCENTILE';

```

### Exercise 3) (1 point)

Now create a file in your working directory called ex3.cql using the Edit-Term. In this file write the commands to insert the following records into table 'Music'...

artistName	albumName	numberSold	cost
<b>Mozart</b>	Greatest Hits	100000	10
<b>Taylor Swift</b>	Fearless	2300000	15
<b>Black Sabbath</b>	Paranoid	534000	12
<b>Katy Perry</b>	Prism	800000	16
<b>Katy Perry</b>	Teenage Dream	750000	14

- a) Execute ex3.cql. Provide the content of this file as the result of this exercise.

```

[hadoop@ip-172-31-44-252 ~]$ cat ex3.cql
INSERT INTO MUSIC (artistname,albumname,numbersold,cost) values ('Mozart','Greatest Hits',100000,10);
INSERT INTO MUSIC (artistname,albumname,numbersold,cost) values ('Taylor Swift','Fearless',2300000,15);
INSERT INTO MUSIC (artistname,albumname,numbersold,cost) values ('Black Sabbath','Paranoid',534000,12);
INSERT INTO MUSIC (artistname,albumname,numbersold,cost) values ('Katy Perry','Prism',800000,16);
INSERT INTO MUSIC (artistname,albumname,numbersold,cost) values ('Katy Perry','Teenage Dream',750000,14);[hadoop@ip-172-31-44-252 ~]$

```

- b) Execute the command 'SELECT \* FROM Music;' and provide the output of this command as another result of the exercise.

```

cqlsh:a20506653> SELECT * FROM Music;

artistname | albumname | cost | numbersold
-----+-----+-----+-----
Mozart     | Greatest Hits | 10 | 100000
Black Sabbath | Paranoid | 12 | 534000
Taylor Swift | Fearless | 15 | 2300000
Katy Perry | Teenage Dream | 14 | 750000
Katy Perry | Prism | 16 | 800000

(5 rows)
cqlsh:a20506653>

```



#### Exercise 4) (1 point)

Now create a file in your working directory called ex4.cql using the Edit-Term. In this file write the commands to query and output only Katy Perry songs. Execute ex4.cql. Provide the content of this file and output of executing this file as the result of this exercise.

```
[hadoop@ip-172-31-44-252 ~]$ cat ex4.cql
select * from music where artistname in ('Katy Perry');
```

```
cqlsh:a20506653> source './ex4.cql';

  artistname | albumname      | cost | numbersold
-----+-----+-----+-----
  Katy Perry | Teenage Dream |   14 |    750000
  Katy Perry | Prism         |   16 |    800000
(2 rows)
cqlsh:a20506653> |
```

#### Exercise 5) (1 point)

Now create a file in your working directory called ex5.cql using the Edit-Term. In this file write the commands to query only albums that have sold 700000 copies or more. Execute ex5.cql. Provide the content of this file and the output of executing this file as the result of this exercise.

```
select * from music where numbersold >= 700000 ALLOW FILTERING;
```

```
cqlsh:a20506653> source './ex5.cql';

  artistname | albumname      | cost | numbersold
-----+-----+-----+-----
  Taylor Swift | Fearless      |   15 |    2300000
    Katy Perry | Teenage Dream |   14 |     750000
    Katy Perry | Prism         |   16 |     800000
(3 rows)
cqlsh:a20506653> |
```

### Step A – Start an EMR cluster

Step B - Download the assignment software (mongoex.tar, mongodb-org-4.2.repo) to master node

### Step C – Install assignment software (mongoex.zip, mongodb-org-4.2.repo)

```
[hadoop@ip-172-31-41-149 ~]$ tar -xvf mongoex.tar
./demo1.js
demo1.js
demo2.js
demo3.js
demo4.js
demo5.js
demo6.js
demo7.js
demo8.js
demo9.js
load.js
[hadoop@ip-172-31-41-149 ~]$
```

### Step D – Install and start MongoDB

Enter the following into Init-Term to install MongoDB:

```
sudo yum install -y mongodb-org-4.2.15 mongodb-org-server-4.2.15 mongodb-org-shell-4.2.15 mongodb-org-mongos-4.2.15 mongodb-org-tools-4.2.15
```

```
Running transaction
Installing : mongodb-org-shell-4.2.15-1.amzn1.x86_64 1/5
Installing : mongodb-org-mongos-4.2.15-1.amzn1.x86_64 2/5
Installing : mongodb-org-tools-4.2.15-1.amzn1.x86_64 3/5
Installing : mongodb-org-server-4.2.15-1.amzn1.x86_64 4/5
Installing : mongodb-org-4.2.15-1.amzn1.x86_64 5/5
Verifying : mongodb-org-4.2.15-1.amzn1.x86_64 1/5
Verifying : mongodb-org-server-4.2.15-1.amzn1.x86_64 2/5
Verifying : mongodb-org-tools-4.2.15-1.amzn1.x86_64 3/5
Verifying : mongodb-org-mongos-4.2.15-1.amzn1.x86_64 4/5
Verifying : mongodb-org-shell-4.2.15-1.amzn1.x86_64 5/5

Installed:
  mongodb-org.x86_64 0:4.2.15-1.amzn1                  mongodb-org-server.x86_64 0:4.2.15-1.amzn1

Dependency Installed:
  mongodb-org-mongos.x86_64 0:4.2.15-1.amzn1          mongodb-org-shell.x86_64 0:4.2.15-1.amzn1          mongodb-org-tools.x86_64 0:4.2.15-1.amzn1

Complete!
```

Now enter this into Init-Term to start mongod:

```
sudo systemctl start mongod
```

### Step E – Start the MongoDB Shell (Command Line Interpreter)

Open a second terminal connection to the EMR master node. Going forward we will call this terminal connection: CLI-Term.

You will use this terminal window to start and run the mongodb shell as follows:

```
mongo
```

```
---
Enable MongoDB's free cloud-based monitoring service, which will then receive and display
metrics about your deployment (disk utilization, CPU, operation statistics, etc).

The monitoring data will be available on a MongoDB website with a unique URL accessible to you
and anyone you share the URL with. MongoDB may use this information to make product
improvements and to suggest MongoDB products and deployment options to you.

To enable free monitoring, run the following command: db.enableFreeMonitoring()
To permanently disable this reminder, run the following command: db.disableFreeMonitoring()
---
> |
```

### Step F – Edit mongo query language files

Open a third terminal connection to the EMR master node. Going forward we will call this terminal connection: CLI-Term. You will use this terminal window to run the 'vi' editor to create your Mongo code files.

As an alternative you could edit your MongoDB code files on your PC/MAC and then 'scp' them to the EMR mater node.

### Step G – Setting up the assignment database

Now, in the MongoDB shell, using the CLI-Term, create a database called "assignment" by entering the following into the MongoDB shell:

```
use assignment;
```

```
> use assignment;
switched to db assignment
> |
```

This will set the shell variable 'db' to this new database.

Load a collection called 'unicorns' with sample data by executing the script load.js in the MongoDB shell as follows (don't cut and paste this, type it in manually):

```
load('./load.js');
```

```
> load('./load.js');
true
> |
```

Note, look at the content of the script file (via the other terminal window you have opened to the EC2 instance) to see how each unicorn is described.

Confirm this has all worked by executing the following command in the MongoDB shell:

```
db.unicorns.find();
```

```
db.unicorns.find()
{ "_id" : ObjectId("6387b18ba0f8592c98827582"), "name" : "Horny", "dob" : ISODate("1992-03-13T07:47:00Z"), "loves" : [ "carrot", "papaya" ], "weight" : 600, "gender" : "m",
  "vampires" : 63 }
{ "_id" : ObjectId("6387b18ba0f8592c98827583"), "name" : "Aurora", "dob" : ISODate("1991-01-24T13:00:00Z"), "loves" : [ "carrot", "grape" ], "weight" : 450, "gender" : "f",
  "vampires" : 43 }
{ "_id" : ObjectId("6387b18ba0f8592c98827584"), "name" : "Unicrom", "dob" : ISODate("1973-02-09T22:10:00Z"), "loves" : [ "energon", "redbull" ], "weight" : 984, "gender" :
  "m", "vampires" : 182 }
{ "_id" : ObjectId("6387b18ba0f8592c98827585"), "name" : "Rooodoodles", "dob" : ISODate("1979-08-18T18:44:00Z"), "loves" : [ "apple" ], "weight" : 575, "gender" : "m", "vam
  pires" : 99 }
{ "_id" : ObjectId("6387b18ba0f8592c98827586"), "name" : "Solnara", "dob" : ISODate("1985-07-04T02:01:00Z"), "loves" : [ "apple", "carrot", "chocolate" ], "weight" : 550, "
  gender" : "f", "vampires" : 80 }
{ "_id" : ObjectId("6387b18ba0f8592c98827587"), "name" : "Ayna", "dob" : ISODate("1998-03-07T08:30:00Z"), "loves" : [ "strawberry", "lemon" ], "weight" : 733, "gender" : "f",
  "vampires" : 40 }
{ "_id" : ObjectId("6387b18ba0f8592c98827588"), "name" : "Kenny", "dob" : ISODate("1997-07-01T10:42:00Z"), "loves" : [ "grape", "lemon" ], "weight" : 690, "gender" : "m", "
  vampires" : 39 }
{ "_id" : ObjectId("6387b18ba0f8592c98827589"), "name" : "Raleigh", "dob" : ISODate("2005-05-03T00:57:00Z"), "loves" : [ "apple", "sugar" ], "weight" : 421, "gender" : "m",
  "vampires" : 2 }
{ "_id" : ObjectId("6387b18ba0f8592c9882758a"), "name" : "Leia", "dob" : ISODate("2001-10-08T14:53:00Z"), "loves" : [ "apple", "watermelon" ], "weight" : 601, "gender" : "f",
  "vampires" : 33 }
{ "_id" : ObjectId("6387b18ba0f8592c9882758b"), "name" : "Pilot", "dob" : ISODate("1997-03-01T05:03:00Z"), "loves" : [ "apple", "watermelon" ], "weight" : 650, "gender" : "
  m", "vampires" : 54 }
{ "_id" : ObjectId("6387b18ba0f8592c9882758c"), "name" : "Nimue", "dob" : ISODate("1999-12-20T16:15:00Z"), "loves" : [ "grape", "carrot" ], "weight" : 540, "gender" : "f" }
{ "_id" : ObjectId("6387b18ba0f8592c9882758d"), "name" : "Dunx", "dob" : ISODate("1976-07-18T18:18:00Z"), "loves" : [ "grape", "watermelon" ], "weight" : 704, "gender" : "m",
  "vampires" : 165 }
```

Note, the files named "demo\*.js" (also included in the mongoex.tar file) provide examples of how to operate in the unicorn collection. These are a VERY good idea to review and understand and will present you with information helpful in completing the assignment. Also, try them out by typing something like

```
load('./demo1.js');
```

```
> load('./demo1.js');
true
> |
```

## Exercises:

### Exercise 1) (1 point)

Write a command that finds all unicorns having weight less than 500 pounds. Include the code you executed and some sample output as the result of this exercise. Recall you can place the command, if you choose, into a file, say 'ex1.js' and execute it with the load command as above and similarly for the following exercises.

```
db.unicorns.find({weight:{ $lt:500}})
```

```
> db.unicorns.find({weight:{ $lt:500}})
{ "_id" : ObjectId("6387b18ba0f8592c98827583"), "name" : "Aurora", "dob" : ISODate("1991-01-24T13:00:00Z"), "loves" : [ "carrot", "grape" ], "weight" : 450, "gender" : "F",
  "vampires" : 43 }
{ "_id" : ObjectId("6387b18ba0f8592c98827589"), "name" : "Raleigh", "dob" : ISODate("2005-05-03T00:57:00Z"), "loves" : [ "apple", "sugar" ], "weight" : 421, "gender" : "m",
  "vampires" : 2 }
```

### Exercise 2) (1 point)

Write a command that finds all unicorns who love apples. Hint, search for "apple". Include the code you executed and some sample output as the result of this exercise.

```
db.unicorns.find({loves:'apple'})
```

```
> db.unicorns.find({loves:'apple'})
{ "_id" : ObjectId("6387b18ba0f8592c98827585"), "name" : "Roooooodles", "dob" : ISODate("1979-08-18T18:44:00Z"), "loves" : [ "apple" ], "weight" : 575, "gender" : "m", "vampires" : 99 }
{ "_id" : ObjectId("6387b18ba0f8592c98827586"), "name" : "Solnara", "dob" : ISODate("1985-07-04T02:01:00Z"), "loves" : [ "apple", "carrot", "chocolate" ], "weight" : 550, "gender" : "F", "vampires" : 80 }
{ "_id" : ObjectId("6387b18ba0f8592c98827589"), "name" : "Raleigh", "dob" : ISODate("2005-05-03T00:57:00Z"), "loves" : [ "apple", "sugar" ], "weight" : 421, "gender" : "m", "vampires" : 2 }
{ "_id" : ObjectId("6387b18ba0f8592c9882758a"), "name" : "Leia", "dob" : ISODate("2001-10-08T14:53:00Z"), "loves" : [ "apple", "watermelon" ], "weight" : 601, "gender" : "f", "vampires" : 33 }
{ "_id" : ObjectId("6387b18ba0f8592c9882758b"), "name" : "Pilot", "dob" : ISODate("1997-03-01T05:03:00Z"), "loves" : [ "apple", "watermelon" ], "weight" : 650, "gender" : "m", "vampires" : 54 }
```

### Exercise 3) (1 point)

Write a command that adds a unicorn with the following attributes to the collection. Note dob means "Date of Birth."

Attribute	Value(s)
name	Malini
dob	11/03/2008
loves	pears, grapes
weight	450
gender	F
vampires	23
horns	1

Include the code you executed to insert this unicorn into the collection along with the output of a find command showing it is in the collection.

```
db.unicorns.insertOne({name:'Malini', dob:'2008-03-11',
loves:['pears','grapes'], weight:450, gender:'F', vampires:23, horns:1})
```

```

> db.unicorns.insertOne({name:'Malini', dob:'2008-03-11', loves:['pears','grapes'], weight:450, gender:'F', vampires:23, horns:1})
{
  "acknowledged" : true,
  "insertedId" : ObjectId("6387b3f2a0f8592c9882758f")
}
>

> db.unicorns.find();
{ "_id" : ObjectId("6387b18ba0f8592c98827582"), "name" : "Horny", "dob" : ISODate("1992-03-13T07:47:00Z"), "loves" : [ "carrot", "papaya" ], "weight" : 600, "gender" : "m", "vampires" : 63 }
{ "_id" : ObjectId("6387b18ba0f8592c98827583"), "name" : "Aurora", "dob" : ISODate("1991-01-24T13:00:00Z"), "loves" : [ "carrot", "grape" ], "weight" : 450, "gender" : "f", "vampires" : 43 }
{ "_id" : ObjectId("6387b18ba0f8592c98827584"), "name" : "Unicrom", "dob" : ISODate("1973-02-09T22:10:00Z"), "loves" : [ "energon", "redbull" ], "weight" : 984, "gender" : "m", "vampires" : 182 }
{ "_id" : ObjectId("6387b18ba0f8592c98827585"), "name" : "Roooooodles", "dob" : ISODate("1979-08-18T18:44:00Z"), "loves" : [ "apple" ], "weight" : 575, "gender" : "m", "vampires" : 99 }
{ "_id" : ObjectId("6387b18ba0f8592c98827586"), "name" : "Solnara", "dob" : ISODate("1985-07-04T02:01:00Z"), "loves" : [ "apple", "carrot", "chocolate" ], "weight" : 550, "gender" : "f", "vampires" : 80 }
{ "_id" : ObjectId("6387b18ba0f8592c98827587"), "name" : "Ayna", "dob" : ISODate("1998-03-07T08:30:00Z"), "loves" : [ "strawberry", "lemon" ], "weight" : 733, "gender" : "f", "vampires" : 40 }
{ "_id" : ObjectId("6387b18ba0f8592c98827588"), "name" : "Kenny", "dob" : ISODate("1997-07-01T10:42:00Z"), "loves" : [ "grape", "lemon" ], "weight" : 690, "gender" : "m", "vampires" : 39 }
{ "_id" : ObjectId("6387b18ba0f8592c98827589"), "name" : "Raleigh", "dob" : ISODate("2005-05-03T00:57:00Z"), "loves" : [ "apple", "sugar" ], "weight" : 421, "gender" : "m", "vampires" : 2 }
{ "_id" : ObjectId("6387b18ba0f8592c9882758a"), "name" : "Leia", "dob" : ISODate("2001-10-08T14:53:00Z"), "loves" : [ "apple", "watermelon" ], "weight" : 601, "gender" : "f", "vampires" : 33 }
{ "_id" : ObjectId("6387b18ba0f8592c9882758b"), "name" : "Pilot", "dob" : ISODate("1997-03-01T05:03:00Z"), "loves" : [ "apple", "watermelon" ], "weight" : 650, "gender" : "m", "vampires" : 54 }
{ "_id" : ObjectId("6387b18ba0f8592c9882758c"), "name" : "Nimue", "dob" : ISODate("1999-12-20T16:15:00Z"), "loves" : [ "grape", "carrot" ], "weight" : 540, "gender" : "f" }
{ "_id" : ObjectId("6387b18ba0f8592c9882758d"), "name" : "Dunx", "dob" : ISODate("1976-07-18T18:18:00Z"), "loves" : [ "grape", "watermelon" ], "weight" : 704, "gender" : "m", "vampires" : 165 }
{ "_id" : ObjectId("6387b1f3a0f8592c9882758e"), "name" : "Leto", "gender" : "m", "home" : "Arrakeen", "worm" : false }
{ "_id" : ObjectId("6387b3f2a0f8592c9882758f"), "name" : "Malini", "dob" : "2008-03-11", "loves" : [ "pears", "grapes" ], "weight" : 450, "gender" : "F", "vampires" : 23, "horns" : 1 }

```

#### Exercise 4) (1 point)

Write a command that updates the above record to add apricots to the list of things Malini loves. Include the code you executed and some sample output showing the addition.

```

db.unicorns.updateOne({loves:['pears','grapes']},{ $set:{loves:['pears','grapes','apricots']}})

```

```

> db.unicorns.updateOne({loves:['pears','grapes']},{ $set:{loves:['pears','grapes','apricots']}})
{ "acknowledged" : true, "matchedCount" : 1, "modifiedCount" : 1 }
>

> db.unicorns.find();
{ "_id" : ObjectId("6387b18ba0f8592c98827582"), "name" : "Horny", "dob" : ISODate("1992-03-13T07:47:00Z"), "loves" : [ "carrot", "papaya" ], "weight" : 600, "gender" : "m", "vampires" : 63 }
{ "_id" : ObjectId("6387b18ba0f8592c98827583"), "name" : "Aurora", "dob" : ISODate("1991-01-24T13:00:00Z"), "loves" : [ "carrot", "grape" ], "weight" : 450, "gender" : "f", "vampires" : 43 }
{ "_id" : ObjectId("6387b18ba0f8592c98827584"), "name" : "Unicrom", "dob" : ISODate("1973-02-09T22:10:00Z"), "loves" : [ "energon", "redbull" ], "weight" : 984, "gender" : "m", "vampires" : 182 }
{ "_id" : ObjectId("6387b18ba0f8592c98827585"), "name" : "Roooooodles", "dob" : ISODate("1979-08-18T18:44:00Z"), "loves" : [ "apple" ], "weight" : 575, "gender" : "m", "vampires" : 99 }
{ "_id" : ObjectId("6387b18ba0f8592c98827586"), "name" : "Solnara", "dob" : ISODate("1985-07-04T02:01:00Z"), "loves" : [ "apple", "carrot", "chocolate" ], "weight" : 550, "gender" : "f", "vampires" : 80 }
{ "_id" : ObjectId("6387b18ba0f8592c98827587"), "name" : "Ayna", "dob" : ISODate("1998-03-07T08:30:00Z"), "loves" : [ "strawberry", "lemon" ], "weight" : 733, "gender" : "f", "vampires" : 40 }
{ "_id" : ObjectId("6387b18ba0f8592c98827588"), "name" : "Kenny", "dob" : ISODate("1997-07-01T10:42:00Z"), "loves" : [ "grape", "lemon" ], "weight" : 690, "gender" : "m", "vampires" : 39 }
{ "_id" : ObjectId("6387b18ba0f8592c98827589"), "name" : "Raleigh", "dob" : ISODate("2005-05-03T00:57:00Z"), "loves" : [ "apple", "sugar" ], "weight" : 421, "gender" : "m", "vampires" : 2 }
{ "_id" : ObjectId("6387b18ba0f8592c9882758a"), "name" : "Leia", "dob" : ISODate("2001-10-08T14:53:00Z"), "loves" : [ "apple", "watermelon" ], "weight" : 601, "gender" : "f", "vampires" : 33 }
{ "_id" : ObjectId("6387b18ba0f8592c9882758b"), "name" : "Pilot", "dob" : ISODate("1997-03-01T05:03:00Z"), "loves" : [ "apple", "watermelon" ], "weight" : 650, "gender" : "m", "vampires" : 54 }
{ "_id" : ObjectId("6387b18ba0f8592c9882758c"), "name" : "Nimue", "dob" : ISODate("1999-12-20T16:15:00Z"), "loves" : [ "grape", "carrot" ], "weight" : 540, "gender" : "f" }
{ "_id" : ObjectId("6387b18ba0f8592c9882758d"), "name" : "Dunx", "dob" : ISODate("1976-07-18T18:18:00Z"), "loves" : [ "grape", "watermelon" ], "weight" : 704, "gender" : "m", "vampires" : 165 }
{ "_id" : ObjectId("6387b1f3a0f8592c9882758e"), "name" : "Leto", "gender" : "m", "home" : "Arrakeen", "worm" : false }
{ "_id" : ObjectId("6387b3f2a0f8592c9882758f"), "name" : "Malini", "dob" : "2008-03-11", "loves" : [ "pears", "grapes", "apricots" ], "weight" : 450, "gender" : "F", "vampires" : 23, "horns" : 1 }

```

#### Exercise 5) (1 point)

Write a command that deletes all unicorns with weight more than 600 pounds. Include the code you executed and some sample output as the result of this exercise.

```

db.unicorns.deleteMany({weight:{ $gt:600}})

```

```

> db.unicorns.deleteMany({weight:{ $gt:600}})
{ "acknowledged" : true, "deletedCount" : 6 }
>

```

```
> db.unicorns.find();
{ "_id" : ObjectId("6387b18ba0f8592c98827582"), "name" : "Horny", "dob" : ISODate("1992-03-13T07:47:00Z"), "loves" : [ "carrot", "papaya" ], "weight" : 600, "gender" : "m",
  "vampires" : 63 }
{ "_id" : ObjectId("6387b18ba0f8592c98827583"), "name" : "Aurora", "dob" : ISODate("1991-01-24T13:00:00Z"), "loves" : [ "carrot", "grape" ], "weight" : 450, "gender" : "f",
  "vampires" : 43 }
{ "_id" : ObjectId("6387b18ba0f8592c98827585"), "name" : "Rooooooodles", "dob" : ISODate("1979-08-18T18:44:00Z"), "loves" : [ "apple" ], "weight" : 575, "gender" : "m", "vampires" : 99 }
{ "_id" : ObjectId("6387b18ba0f8592c98827586"), "name" : "Solnara", "dob" : ISODate("1985-07-04T02:01:00Z"), "loves" : [ "apple", "carrot", "chocolate" ], "weight" : 550, "gender" : "f", "vampires" : 80 }
{ "_id" : ObjectId("6387b18ba0f8592c98827589"), "name" : "Raleigh", "dob" : ISODate("2005-05-03T00:57:00Z"), "loves" : [ "apple", "sugar" ], "weight" : 421, "gender" : "m", "vampires" : 2 }
{ "_id" : ObjectId("6387b18ba0f8592c9882758c"), "name" : "Nimue", "dob" : ISODate("1999-12-20T16:15:00Z"), "loves" : [ "grape", "carrot" ], "weight" : 540, "gender" : "f" }
{ "_id" : ObjectId("6387b1f3a0f8592c9882758e"), "name" : "Leto", "gender" : "m", "home" : "Arrakeen", "worm" : false }
{ "_id" : ObjectId("6387b3f2a0f8592c9882758f"), "name" : "Malini", "dob" : "2008-03-11", "loves" : [ "pears", "grapes", "apricots" ], "weight" : 450, "gender" : "F", "vampires" : 23, "horns" : 1 }
>
```